

# Clustering

Filipe Cabral Pinto


# Key concept

- Clustering is an unsupervised machine learning topic intending to find data patterns having no target to predict
- Clustering aims at grouping similar data points in large datasets
  - A cluster should group similar observations
  - Observations with distinguishable characteristics should belong to different clusters
- Similarity can be defined in different ways

# Clustering few examples

- Customer segmentation
- Network anomalies
- Fraud detection
- Document Analysis
- Streaming Services
- Sports science

# Grouping data points



# Clustering approaches

## Centroid-based models

**Data points are grouped based on the proximity to the centroids.**

**k-means is the most well-known centroid-based clustering algorithm.**

## Density-based models




**Data points concentrate over contiguous regions, and empty or sparse areas separate clusters.**

**DBSCAN is probably the most widely-used algorithm.**

# K-Means algorithm

- K-Means is an unsupervised algorithm aiming to group unlabelled data points into K distinct clusters.
  - A data point belongs to one and only one cluster.
- K-Means finds similar data points and group them together.
- K-Means uses Euclidean distances to measure the proximity.
  - It allocates each data point dynamically to the nearest cluster centred in one of the K centroids identified.

# K-Means




# Inputs and outputs

## Inputs

Dataset:  $S_n = \{x^{(i)}, i = 1 \dots n\}$

Number of clusters:  $K$

**Note:**  $x^{(i)}$  =  $i^{\text{th}}$  data point of the dataset




## Outputs

Set of partitions:  $C_1, C_2, \dots, C_k$

$$\bigcup_{k=1}^K C_k = S_n$$

$$C_i \cap C_j = \emptyset, i \neq j$$





**Note:**  $C_k$  includes a set of data points





# How to group?

$K = 2 \Rightarrow 2$  Clusters

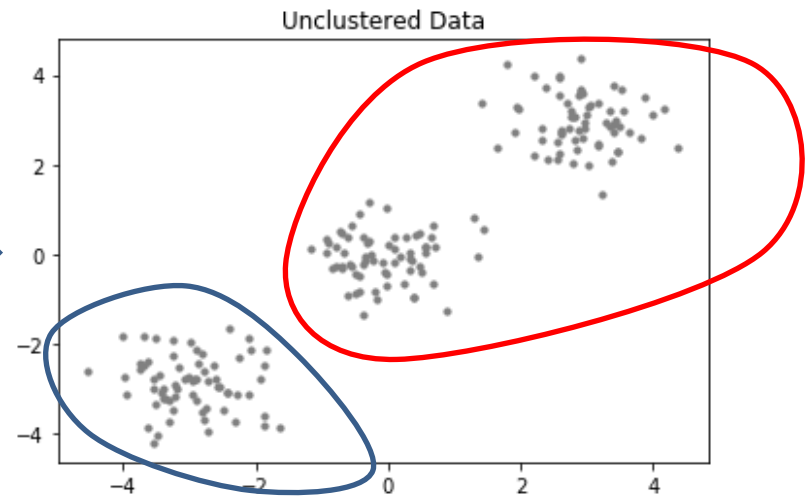




# Total cost or loss function

Cost of all  
clusters

$$\text{Cost}(C_1, C_2, \dots, C_k) = \sum_{k=1}^K \text{Cost}(C_k)$$

$C_k$  is the  $k^{\text{th}}$  cluster and includes a set of data points



# Cluster cost

$$\text{Cost}(C_k, c^{(k)}) = \sum_{x^{(i)} \in C_k} \text{distance\_metric}(c^{(k)}, x^{(i)})$$

$$\text{Cost}(C_k, c^{(k)}) = \sum_{x^{(i)} \in C_k} \|c^{(k)} - x^{(i)}\|^2 = \sum_{x^{(i)} \in C_k} (c^{(k)} - x^{(i)})^2$$

$C_k$  is the  $k^{\text{th}}$  cluster and includes a set of data points


$c^{(k)}$  is the representative centroid of cluster  $C_k$

$x^{(i)}$  is the  $i^{\text{th}}$  data point of the dataset

**Note:**  $\|c - x\| = \sqrt{(c_1 - x_1)^2 + (c_2 - x_2)^2 + \dots + (c_m - x_m)^2}$

$m$  is the number of features

Cost of the  $k$   
cluster



# K-Means best partitions iterative steps

$$\text{Cost}(C_1, C_2, \dots, C_k, c^{(1)}, \dots, c^{(1)}) = \sum_{k=1}^K \sum_{x^{(i)} \in C_k} \|c^{(k)} - x^{(i)}\|^2$$

**Objective:** To minimize the **total cost**

Given the representative centroids -  $c^{(1)}, c^{(2)}, \dots, c^{(K)}$  - allocate each  $x$  to the closest  $c$

$$\text{Cost}(c^{(1)}, c^{(2)}, \dots, c^{(K)}) = \sum_{i=1}^n \min_{k=1 \dots K} (\|c^{(k)} - x^{(i)}\|^2)$$

Use Euclidean distance to select the closest  $c$

Given the Clusters -  $C_1, C_2, \dots, C_K$  - select the best representatives centroids  $c$

That's the tricky part!

$$\text{Cost}(C_1, C_2, \dots, C_k) = \min_{c^{(1)}, c^{(2)}, \dots, c^{(K)}} \sum_{k=1}^K \sum_{x^{(i)} \in C_k} \|c^{(k)} - x^{(i)}\|^2$$


Use the mean to find the centroids

# Euclidean distance

The most widely used distance function is the Euclidean distance

$$\text{distance}(A, B) = \sqrt{\sum_{i=1}^m (a_i - b_i)^2}$$

$A = (a_1, a_2, \dots, a_m)$  and  $B = (b_1, b_2, \dots, b_m)$



$$\text{distance}(A, B) = \sqrt{\sum_{i=1}^m (a_i - b_i)^2} = \sqrt{(a_2 - b_2)^2 + (a_1 - b_1)^2}$$

**Remember  
Pythagoras Theorem**

# The mean from K-Means

$$c^{(j)} = \frac{\sum_{x^{(i)} \in C_j} (x^{(i)})}{|C_j|}$$

K-Means uses the mean of each cluster to find the new representative centroids

$$\text{Cost}(C_1, C_2, \dots, C_K) = \min_{c^{(1)}, c^{(2)}, \dots, c^{(K)}} \sum_{k=1}^K \sum_{x^{(i)} \in C_k} \|c^{(k)} - x^{(i)}\|^2$$

Partial derivative

$$\frac{\partial C}{\partial c^{(j)}} = \frac{\partial}{\partial c^{(j)}} \left( \sum_{k=1}^K \sum_{x^{(i)} \in C_k} \|c^{(k)} - x^{(i)}\|^2 \right) \Rightarrow \frac{\partial C}{\partial c^{(j)}} = \frac{\partial}{\partial c^{(j)}} \left( \sum_{k=1}^K \sum_{x^{(i)} \in C_k} (c^{(k)} - x^{(i)})^2 \right)$$

$$\frac{\partial C}{\partial c^{(j)}} = \sum_{x^{(i)} \in C_j} \frac{\partial}{\partial c^{(j)}} (c^{(j)} - x^{(i)})^2 \Rightarrow \frac{\partial C}{\partial c^{(j)}} = \sum_{x^{(i)} \in C_j} 2 \cdot (c^{(j)} - x^{(i)})$$

# The mean from K-Means

$$c^{(j)} = \frac{\sum_{x^{(i)} \in C_j} (x^{(i)})}{|C_j|}$$

K-Means uses the mean of each cluster to find the new representative centroids

$$\sum_{x^{(i)} \in C_j} 2 \cdot (c^{(j)} - x^{(i)}) = 0 \quad \Rightarrow \quad \sum_{x^{(i)} \in C_j} (c^{(j)} - x^{(i)}) = 0$$

... continuation

$$\sum_{x^{(i)} \in C_j} (c^{(j)}) - \sum_{x^{(i)} \in C_j} (x^{(i)}) = 0 \quad \Rightarrow \quad c^{(j)} \cdot \sum_{x^{(i)} \in C_j} (1) - \sum_{x^{(i)} \in C_j} (x^{(i)}) = 0$$

$$c^{(j)} \cdot |C_j| = \sum_{x^{(i)} \in C_j} (x^{(i)}) = 0 \quad \Rightarrow \quad c^{(j)} = \frac{\sum_{x^{(i)} \in C_j} (x^{(i)})}{|C_j|}$$

The new representative centroid of a specific cluster is the mean of all data points in that cluster

**Note:**  $|C_j|$  is the cardinality of the set of  $C_j$

# K-Means steps

Consider K centroids randomly

It is typical to select K random points from the dataset

For each data point, check the distances to the centroids

Use Euclidean distance

Associate each point with the closest centroid

Centroids are the cluster representatives


Reassign the centroids positions

Set centroids to the mean of each cluster

Repeat the process until there are no more changes in the clusters.




# Dataset



	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

**K = 2**


# Target



	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

**K = 2**

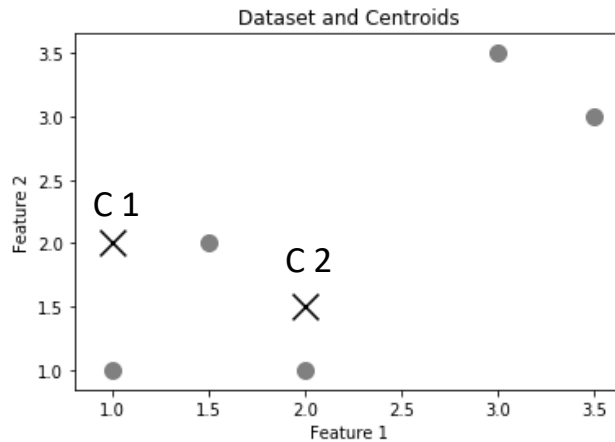
# Point 1



	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0


# Centroids

There are different initialization methods with varying levels of complexity.



C 1	Feature 1	Feature 2
C 1	1.0	2.0
C 2	2.0	1.5

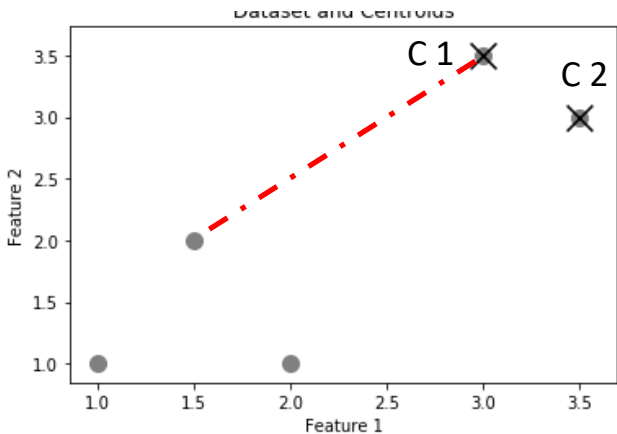

Random selection from the data space is an inefficient possibility



	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

Random selection from the dataset is the common one

# Distance between points and centroids




Different metrics can be applied

	Feature 1	Feature 2
C 1	1.0	2.0
C 2	2.0	1.5

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0


# Dataset and centroids



	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

# Point 0 to C1




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_0, C_1) = \sqrt{(1 - 3)^2 + (1 - 3.5)^2} = 3.20$$

# Point 0 to C2




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_0, C_2) = \sqrt{(1 - 3.5)^2 + (1 - 3)^2} = 3.20$$



# Grouping Point 0




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

Point 0 is the same distance from C1 and C2,  
so assume it belongs to Cluster 1

# Point 1 to C1

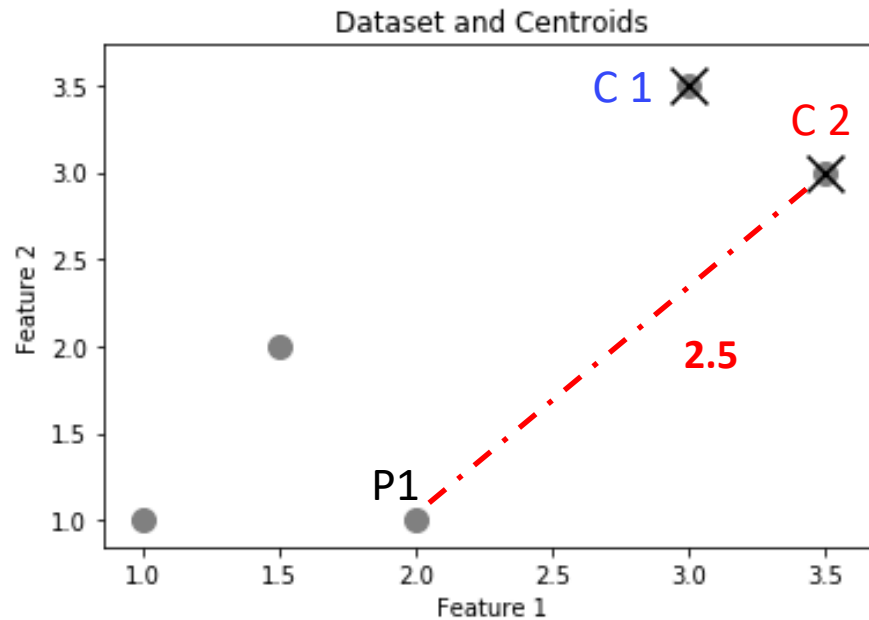


	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_1, C_1) = \sqrt{(2 - 3)^2 + (1 - 3.5)^2} = 2.69$$

# Point 1 to C2

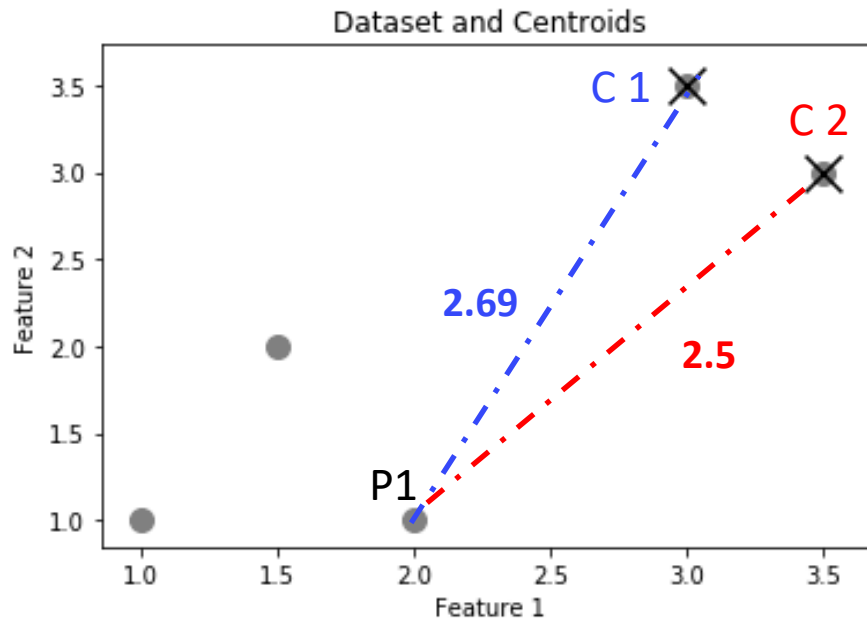


	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_1, C_2) = \sqrt{(2 - 3.5)^2 + (1 - 3)^2} = 2.5$$

# Grouping Point 1




	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	●
2	1.5	2.0	
3	3.0	3.5	
4	3.5	3.0	

	Feature 1	Feature 2	
C 1	3.0	3.5	●
C 2	3.5	3.0	●

Point 1 is closer to C2 so it belongs to Cluster 2

# Point 2 to C1




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_2, C_1) = \sqrt{(1.5 - 3)^2 + (2 - 3.5)^2} = 2.12$$

# Point 2 to C2

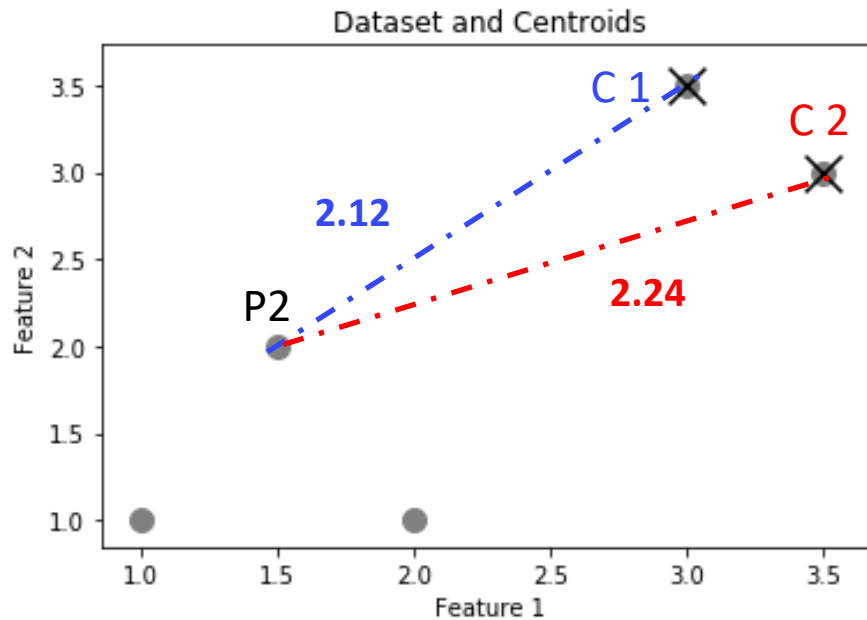


	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_2, C_2) = \sqrt{(1.5 - 3.5)^2 + (2 - 3)^2} = 2.24$$

# Grouping Point 2




	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	●
2	1.5	2.0	●
3	3.0	3.5	
4	3.5	3.0	

	Feature 1	Feature 2	
C 1	3.0	3.5	●
C 2	3.5	3.0	●

Point 2 is closer to C1 so it belongs to Cluster 1



# Point 3 to C1




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_3, C_1) = \sqrt{(3 - 3)^2 + (3.5 - 3.5)^2} = 0$$



# Point 3 to C2




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_3, C_2) = \sqrt{(3 - 3.5)^2 + (3.5 - 3)^2} = 0.71$$

# Grouping Point 3




	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	●
2	1.5	2.0	●
3	3.0	3.5	●
4	3.5	3.0	

	Feature 1	Feature 2	
C 1	3.0	3.5	●
C 2	3.5	3.0	●

Point 3 is closer to C1 so it belongs to Cluster 1



# Point 4 to C1




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_4, C_1) = \sqrt{(3.5 - 3)^2 + (3 - 3.5)^2} = 0.71$$

# Point 4 to C2




	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0

$$D(P_4, C_2) = \sqrt{(3.5 - 3.5)^2 + (3 - 3)^2} = 0$$

# Grouping Point 4




	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	●
2	1.5	2.0	●
3	3.0	3.5	●
4	3.5	3.0	●

	Feature 1	Feature 2	
C 1	3.0	3.5	●
C 2	3.5	3.0	●

Point 4 is closer to C2 so it belongs to Cluster 2



# Data clustering after 1<sup>st</sup> iteration



Remember:  $K = 2$

	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	●
2	1.5	2.0	●
3	3.0	3.5	●
4	3.5	3.0	●

	Feature 1	Feature 2	
C 1	3.0	3.5	●
C 2	3.5	3.0	●

# Centroids reassignment

	Feature 1	Feature 2	
C 1	3.0	3.5	●
C 2	3.5	3.0	

	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	●
2	1.5	2.0	●
3	3.0	3.5	●
4	3.5	3.0	●

	Feature 1	Feature 2	
C 1	3.0	3.5	
C 2	3.5	3.0	●

C 1

C 2

	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	
2	1.5	2.0	●
3	3.0	3.5	●
4	3.5	3.0	

	Feature 1	Feature 2	
0	1.0	1.0	
1	2.0	1.0	●
2	1.5	2.0	
3	3.0	3.5	
4	3.5	3.0	●

# New centroids

	Feature 1	Feature 2	
0	1.0	1.0	●
1	2.0	1.0	
2	1.5	2.0	●
3	3.0	3.5	●
4	3.5	3.0	



$$\text{New C 1\_F1} = \frac{1 + 1.5 + 3}{3} = 1.83$$

$$\text{New C 1\_F2} = \frac{1 + 2 + 3.5}{3} = 2.17$$

	Feature 1	Feature 2	
0	1.0	1.0	
1	2.0	1.0	●
2	1.5	2.0	
3	3.0	3.5	
4	3.5	3.0	●




$$\text{New C 2\_F1} = \frac{2 + 3.5}{2} = 2.75$$

$$\text{New C 2\_F2} = \frac{1 + 3}{2} = 2$$



# Dataset and new centroids





	Feature 1	Feature 2
0	1.0	1.0
1	2.0	1.0
2	1.5	2.0
3	3.0	3.5
4	3.5	3.0

	Feature 1	Feature 2
C 1	1.83	2.17
C 2	2.75	2.00

**Repeat the whole process with the new centroids!**

# The whole process: iteration 1

## 1<sup>st</sup> Iteration




	Feature 1	Feature 2
C 1	3.0	3.5
C 2	3.5	3.0



Random  
centroids from  
the dataset



# The whole process: iteration 2

## 2<sup>nd</sup> Iteration



	Feature 1	Feature 2
C 1	1.83	2.17
C 2	2.75	2.00


Centroids  
calculated in the  
first iteration



Check



# The whole process: iteration 3

## 3<sup>rd</sup> Iteration






	Feature 1	Feature 2
<b>C 1</b>	1.50	1.33
<b>C 2</b>	3.25	3.25

Centroids  
calculated in the  
second iteration





No reassignments!  
Centroids do not  
change anymore!

# The whole process: stopping criteria




1 – There are no reassignments, so centroids do not change their values anymore

2 - The defined maximum number of iterations has been reached



# Final total cost: compute it

$$\text{Cost}(C_1, C_2, \dots, C_k, c^{(1)}, \dots, c^{(k)}) = \sum_{k=1}^K \sum_{x^{(i)} \in C_k} \|c^{(k)} - x^{(i)}\|^2$$



The K-Means cost monotonically decreases

The K-Means algorithm converges to a local minimum

The cost depends on the representative centroids initialization


Sklearn runs by default 10 times using different centroids and selects the one with the lower final total cost, and defines 300 as the default maximum number of iterations

# Plotting the dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
```

```
descriptive = {'Feature 1':[1,2,1.5,3,3.5], 'Feature 2':[1,1,2,3.5,3]}
dataset_features = pd.DataFrame(descriptive)
```

```
plt.scatter(dataset_features['Feature 1'].values,
            dataset_features['Feature 2'].values,
            s=100, c='grey', label='Data')
plt.title('Dataset')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```



# Plotting the dataset and centroids

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans


descriptive = {'Feature 1':[1,2,1.5,3,3.5], 'Feature 2':[1,1,2,3.5,3]}
dataset_features = pd.DataFrame(descriptive)
centroids = np.array([[3,3.5], [3.5,3]])

indexes = ['C1', 'C2']
column_names = ['Feature 1', 'Feature 2']
dataset_centroids = pd.DataFrame(centroids,
                                  index=indexes, columns=column_names)

plt.scatter(dataset_features['Feature 1'].values,
            dataset_features['Feature 2'].values,
            s=100, c='grey', label='Data')

plt.title('Dataset and Centroids')
plt.scatter(dataset_centroids['Feature 1'].values,
            dataset_centroids['Feature 2'].values,
            s=250, marker='x', c='black')

plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```





# Computing and plotting the clusters

```
from sklearn.cluster import KMeans
import numpy as np

k_means = KMeans(n_clusters = 2,
                  random_state = 0, n_init = 1,
                  init = dataset_centroids.values).fit(dataset_features.values[:, :2])

dataset_features['Cluster'] = k_means.labels_


group_1 = dataset_features.loc[dataset_features['Cluster'] == 0]
group_2 = dataset_features.loc[dataset_features['Cluster'] == 1]

plt.scatter(group_1['Feature 1'].values, group_1['Feature 2'].values ,
            s=100, c='b', label='Group 1')
plt.scatter(group_2['Feature 1'].values, group_2['Feature 2'].values ,
            s=100, c='r', label='Group 2')
plt.scatter(k_means.cluster_centers_[ :,0], k_means.cluster_centers_[ :,1],
            s=250, marker = 'x', c='black')

plt.legend()
plt.title('Clustered Data')
plt.show()
```


init sets the initial representative centroids

Sklearn provides the KMeans class



# Final values

```
print('K-Means labels\n',k_means.labels_)
print('K-Means cluster centers\n',k_means.cluster_centers_)
print('Cost\n', k_means.inertia_)
print('Iterations\n', k_means.n_iter_)
```



K-Means labels

[0 0 0 1 1]

K-Means cluster centers

[[1.5            1.33333333]  
[3.25           3.25        ]]

Iterations

3

Cost

1.4166666666666667